

Cited but Not Correct: Auditing Evidence-Grounded LLM Configuration Advice against Equally Informed Baselines

Anonymous ACL submission

Abstract

Language models increasingly recommend actions and justify them with citations to evidence in the prompt. Such a rationale makes three promises: the action is good, its claims about the evidence are true, and the decision rests on what is cited. Evaluations rarely test all three, or compare against what the evidence alone would recommend. Our audit protocol gives an equally informed non-LLM selector the same frozen evidence, ablates and deletes evidence, verifies every quantitative claim in an explanation against it, and scores decisions on independently labeled held-out outcomes. We apply it to configuration advice on real Linux kernel interventions with GroundKern, a testbed separating evidence, proposals and enforcement, and ten LLMs from 7B to frontier scale. The promises come apart. Evidence is necessary: without it, nine models mostly choose worse. It is also sufficient: with it, nine reproduce a table-lookup selector and none improves on it. Citations are valid but one-sided: 252 of 260 cited runs belong to the chosen configuration. Explanations are cited but not correct: 41 of 97 audited claims misstate the evidence, mostly alongside a correct decision, and deletion reveals citation-specific reliance for only one model. Prompts, responses, annotations, code and logs are provided as supplementary material.

1 Introduction

Language models are increasingly asked not only to answer questions but to recommend actions, and to justify each recommendation with citations to evidence supplied in the prompt. A rationale of this kind makes three distinct promises: that the recommended action is good, that what the rationale says about the evidence is true, and that the cited evidence is what the recommendation rests on. Evaluations usually check at most one of them. Work on attribution asks whether generated statements are

Choice: s50000_aall (the best configuration) Cited: training-table, five retrieval runs (all valid IDs)
<i>“... For s50000_aall, the P95 latencies across 10 runs are consistently low: 126.36, 140.554, 94.81, ... 87.347 μs. The maximum among these is 183.217 μs, <u>which is lower than the lowest P95 observed for any other config...</u>”</i>
Check: every quoted value is correct; the underlined comparison is false, since s50000_aone has a run at 85.033 μ s.

Figure 1: A real full-context response from DeepSeek-V3.2 (seed 4089), abridged. The decision is correct and every citation is valid, yet the explanation misstates the evidence it cites. Across ten LLMs, 41 of 97 audited claims are misstated in this way (§5).

supported by their sources (Rashkin et al., 2023; Gao et al., 2023b; Min et al., 2023; Liu et al., 2023), typically for question answering or summarization, where no downstream action is scored. Work on LLM decision-making scores the action (Yang et al., 2024; Lao et al., 2024; Giannakouris and Trummer, 2025) but treats the rationale as commentary. Neither asks what the evidence alone would have recommended.

We study these promises in configuration advice, an application where LLMs read manuals and measurements and propose database or kernel settings with natural-language justifications (Trummer, 2022; Lao et al., 2024; Giannakouris and Trummer, 2025; Chen et al., 2024; Lin et al., 2025). The setting is unusually well suited to auditing grounded rationales. The evidence is small, structured and numeric, so most claims about it can be checked exactly rather than by judging entailment. The quality of a decision can be measured on held-out runs, independently of the model and of any mechanism that enforces its advice. And the evidence can be handed unchanged to a conventional selector, which says what the evidence determines without a language model.

We ask three questions. *Added value:* does the

model choose better than a conventional selector that receives the same evidence, or does it only repackage that evidence? *Explanation correctness*: when an explanation cites valid evidence, does it state that evidence correctly? *Reliance*: does the cited evidence actually influence the decision? To answer them we build GroundKern, a testbed that separates evidence (a typed interaction graph and retrieved measurement traces), proposals (an LLM that emits a configuration with citations and an explanation) and enforcement (a runtime that validates, vetoes or stages the proposal). Because the evidence is an explicit, frozen object, it can be given to a non-LLM selector, ablated, or partially deleted; because enforcement is separate, its errors can be counted against labels computed from raw measurements.

Auditing ten LLMs, from 7B open-weight to hosted frontier models, on real Linux timer-slack and CPU-affinity interventions, we find that the three promises come apart. Evidence is necessary: without it, nine models mostly choose a worse configuration, and the local models we audited justify that choice by inventing meanings for opaque configuration names. It is also sufficient: with it, nine models reproduce the choice of a selector that simply reads the training table, and none improves on it. Citations are valid but one-sided and front-loaded: 252 of 260 cited retrieval runs belong to the chosen configuration, and 209 come from the first half of the evidence list. Explanations are cited but not correct (Figure 1): 16 of 28 audited claims from local models and 25 of 69 from hosted models misreport values, ranges or units, attribute a value to the wrong source, or overstate a regularity, and most of these accompany a correct decision. Deleting cited evidence reveals citation-specific reliance for only one model.

Contributions. (i) An audit protocol for evidence-grounded LLM advice that combines a matched-information non-LLM baseline, evidence-source ablations, cited versus matched uncited deletion, claim-level verification of explanations against the supplied evidence, and outcome labels independent of enforcement, with all data roles frozen before test data are opened (§4). (ii) An audit of ten LLMs reported with every denominator, including an analysis of what models cite and a claim-level taxonomy of explanation errors (§5). (iii) The GroundKern testbed, whose calibrated risk gate provides a cautionary result

on selective prediction under shift, with prompts, raw responses, per-claim annotations, code and measurement logs provided as supplementary material (§3).

The task is deliberately small, one workload family on one host, and its default configuration is best, so the audit can reveal failures of added value but not successes; we make no LLM-superiority, deployment-efficacy or production-safety claim (see Limitations).

2 Related Work

Our audit examines a cited rationale at four levels: the decision, the citations, the claims made about the cited evidence, and the reliance of the decision on that evidence. We relate each level to prior work, then turn to LLMs as decision makers and to risk control.

Citation quality and attribution. Retrieval-augmented generation conditions outputs on retrieved documents (Lewis et al., 2020), and attribution research asks whether the resulting text is supported by what it cites. AIS defines attribution to identified sources (Rashkin et al., 2023); Attributed QA and ALCE benchmark attribution and citation quality (Bohnet et al., 2022; Gao et al., 2023b); RARR revises text to restore attribution (Gao et al., 2023a); and generative search engines leave many statements unsupported by the sources they cite (Liu et al., 2023). These studies ask whether each citation supports its statement. We additionally ask which evidence gets cited at all, and find citations that are valid yet one-sided and concentrated at the start of the evidence list (§5.2).

Claim-level accuracy against data. FActScore checks atomic facts against a knowledge source (Min et al., 2023). Data-to-text generation has long documented factual errors in text produced from tables (Wiseman et al., 2017); Thomson and Reiter (2020) annotate every error in generated basketball summaries against the source data, including incorrect numbers, names and words, and summarization work distinguishes intrinsic from extrinsic hallucination (Maynez et al., 2020; Ji et al., 2023). Statements requiring comparison, superlatives or aggregation over a table are hard both to verify and to generate faithfully (Chen et al., 2020b,a). We apply this error-annotation practice to the free-text rationales

of instruction-tuned LLMs, where the supplied evidence is small enough for every claim to be checked exactly, and find the same superlative, range and quantifier errors (§5.3).

Faithfulness and reliance. Faithfulness must be evaluated separately from plausibility (Jacovi and Goldberg, 2020). ERASER measures the comprehensiveness and sufficiency of rationales by removing them (DeYoung et al., 2020); counterfactual and reconstruction tests probe natural-language explanations (Atanasova et al., 2023); and chain-of-thought explanations can omit the factors that drive a prediction (Turpin et al., 2023) or be tested by intervening on the reasoning (Lanham et al., 2023). Self-explanation faithfulness varies across models and tasks (Madsen et al., 2024), and many such tests measure output-level self-consistency rather than faithfulness to internal computation (Parcalabescu and Frank, 2024); ContextCite attributes generations to context by ablation (Cohen-Wang et al., 2024). Our matched deletion is an ERASER-style comprehensiveness test with a same-kind uncited control and, in the terms of Parcalabescu and Frank (2024), a self-consistency test (§5.4).

How models use supplied context. Models may ignore context that conflicts with parametric knowledge (Longpre et al., 2021) or accept coherent counter-evidence while favoring their priors (Xie et al., 2024); irrelevant context distracts them (Shi et al., 2023), and information in the middle of long contexts is used less (Liu et al., 2024). Our evidence ablation examines the complementary case in which the context is decisive and the prior is misleading (§5.1), and our citation analysis adds a position effect in what models choose to cite.

LLMs as decision makers. LLMs can optimize by proposing candidates from past evaluations (Yang et al., 2024), and configuration advisors mine tuning hints from manuals (Trummer, 2022), prune Bayesian-optimization searches with LLM-extracted knowledge (Lao et al., 2024), generate complete configurations (Giannakouris and Trummer, 2025) or tune Linux kernel build options as agents (Chen et al., 2024; Lin et al., 2025); classical tuners use Gaussian processes, reinforcement learning or black-box optimization (Van Aken et al., 2017; Zhang et al., 2019; Herodotou et al., 2020; Snoek et al., 2012; Golovin et al., 2017; Lee et al., 2022). These works score the chosen action end to end, whereas we ask what the model adds over

a selector given the same evidence and whether its stated reasons are correct; we do not compare with these systems head to head. Following arguments against random or standard splits (Gorman and Bedrick, 2019; Søgaard et al., 2021) and for reporting matched budgets (Dodge et al., 2019), our data roles come from disjoint operating conditions and every baseline has an explicit budget.

Uncertainty and risk control. Self-consistency aggregates sampled outputs (Wang et al., 2023); LLMs can partly predict their own correctness (Kadavath et al., 2022), verbalized confidence can be better calibrated than token probabilities for RLHF models (Tian et al., 2023) yet remains overconfident (Xiong et al., 2024), and semantic entropy measures uncertainty over meanings (Kuhn et al., 2023). Selective classification abstains to meet a target risk (Geifman and El-Yaniv, 2017), and conformal methods give guarantees for language-model outputs (Quach et al., 2024), their factuality (Mohri and Hashimoto, 2024) and monotone losses (Angelopoulos et al., 2024), with extensions for non-exchangeable or drifting data (Farinhas et al., 2024; Barber et al., 2023; Gibbs and Candès, 2021); none implies that a window reset after a detected change (Bifet and Gavalda, 2007) restores a fixed risk level. Our rank gate is a restricted class-conditional instance of this family, and we report where its assumptions fail (§5.5).

3 Testbed: GroundKern

GroundKern is an evidence-grounded generation pipeline for system configuration that we use as an instrument (Figure 2). Four properties matter for the audit. The evidence is *addressable* (every citation resolves to one item), *portable* (a non-LLM selector can receive the same object) and *editable* (kinds or items can be removed), and enforcement is *independent* of generation, so its errors can be counted against labels the generator never sees.

Evidence as a structured context. The store holds typed items, each with a stable ID, a kind and a provenance record: summary tables, individual measurement traces, and edges of a knob-interaction graph. The graph is a small knowledge representation with a strict admission rule: for a lower-is-better loss L and baseline values a_0, b_0 , an edge requires the contrast $I_{ab} = L(a, b) - L(a, b_0) - L(a_0, b) + L(a_0, b_0)$ from independently measured configurations, with a bootstrap interval




Figure 2: GroundKern decision path. Evidence, generation and enforcement are separate components: the frozen evidence object can be given to a non-LLM selector, ablated or partially deleted, citations can be resolved against its item IDs, and gate errors can be counted against independently defined labels. The REST runtime is a dry run; a separate per-thread actuator applies, verifies and restores timer slack and affinity in owned child processes.

over runs that excludes zero, and records its scope, values, measurement IDs and interval. Edges describe empirical interactions, not causal couplings (Appendix A). For the audit, the evidence is frozen and serialized as a JSON list that is identical for every advisor.

Generation. The reasoner retrieves typed neighbors and lexically matched trace text (Lewis et al., 2020) and asks a model for a complete proposal, a set of knob values, with cited item IDs and an explanation. In deployment it draws $k = 3$ samples (Wang et al., 2023) and returns a modal complete proposal actually present in a response, rather than splicing per-knob winners into an unseen combination, scores each member by sample disagreement u (a heuristic, not a calibrated probability; Appendix A), and abstains when no model responds or the context budget would force evidence to be truncated. The audit in §4 calls each model directly on the frozen evidence, one response per seed, so that decisions, citations and explanations can be compared across models without the aggregation step.

Enforcement. The runtime treats a proposal as a unit: it validates every member (allowlisted knobs, bounds, joint constraints, no duplicates) and, if any member score satisfies $u \geq \tau$, vetoes all of them; otherwise the proposal enters a staged rollout. The

Table 1: Audit levels, in the order of §5; all use the same frozen evidence.

Level	Manipulation	Measure
Decision	matched non-LLM selector; evidence ablation	held-out P95 change vs. control
Citation	— (observational)	kind, configuration, position
Claim	— (observational)	labels against the evidence
Reliance	cited vs. uncited matched deletion	decision change

REST runtime is a dry run that never writes kernel settings (Appendix A). The threshold τ comes from a rank gate, a selective-prediction rule (Geifman and El-Yaniv, 2017) calibrated on unsafe examples only. Let $Y = 1$ denote an unsafe outcome defined from raw observations, never from whether a veto or rollback occurred.

Proposition 1 (Restricted rank gate). *Fix a score function before calibration. Suppose n unsafe calibration scores and the next unsafe test score are jointly exchangeable. Set $j = \lfloor \alpha(n + 1) \rfloor$ for $0 < \alpha < 1$. If $j = 0$, veto all; otherwise let τ be the j th smallest unsafe calibration score and accept only when $u < \tau$. Then the unsafe-class acceptance probability, marginal over calibration and test scores, is at most $j/(n + 1) \leq \alpha$.*

Proof. By exchangeability the test score’s rank among the $n + 1$ scores is uniform; it falls below the j th calibration score only in the first j ranks, and ties only reduce this. $\square$

The bound controls misses among unsafe actions, $\Pr(A = 1 \mid Y = 1)$, not the unsafe fraction among accepted actions; a cost-selected threshold is capped at the rank threshold, and no unsafe calibration example implies veto-all. The gate accepts any fixed score; §5.5 evaluates it with a measured risk score rather than u . An expected-cost identity, a discussion of drift and the bounded per-thread actuator used in §5.5 are in Appendix A.

4 Audit Protocol

The protocol treats configuration advice as a grounded generation task and audits its output at four levels, each with its own manipulation and measure (Table 1).

Task. The input is a list of four allowed configuration labels and a JSON evidence list; the output is one label, up to six cited evidence IDs and a free-text explanation. The labels name

per-thread settings: timer slack of $50\mu\text{s}$ or 1 ms crossed with affinity to all four CPUs or to CPU 0 (s50000_aall, s50000_aone, s1000000_aall, s1000000_aone). They are opaque by design: nothing in a label says which setting is better. The objective is the 95th-percentile (P95) wake-up latency; s50000_aall is the control. Runs execute in fresh child processes on one Linux host (Appendix B).

Data roles and frozen evidence. Training, retrieval, calibration and test data come from disjoint wake-up periods of 2, 3, 4 and $\{5, 8\}$ ms (Søgaard et al., 2021). Each period has ten randomized blocks of the four configurations with 40 recorded events per run: 200 runs and 8,000 events, collected under a plan saved in advance. The evidence contains three kinds of items: the training table (mean training-run P95 per configuration), an interaction-graph item (the training contrast, $-18.6\mu\text{s}$, whose prespecified interval includes zero, so the graph has no edge) and 40 retrieval items, one per retrieval run, listed in block order. Selections and thresholds are frozen, with hashes, before calibration and test data are opened; test outcomes never enter a prompt.

Decisions. The equally informed baseline is a table selector that picks the minimum of the training table; a saturated two-factor model of the same means is equivalent. It uses the same training runs as the LLMs and no online trials; matched-budget bandit, RL and BO optimizers (Auer et al., 2002; Snoek et al., 2012) are further references (Appendix B). Ten LLMs receive the same system prompt, evidence and output schema (Appendix C) at temperature 0 with five seeds: Qwen2.5 7B and 14B (Qwen Team, 2024), Llama-3.1-8B (Grattafiori et al., 2024) and Phi-4 (Abdin et al., 2024) locally as hashed 4-bit GGUF files in llama.cpp¹, and Llama-3.3-70B, Qwen3-235B-A22B (Yang et al., 2025), DeepSeek-V3.2 (DeepSeek-AI, 2025), Nemotron-3-Ultra-550B (NVIDIA, 2026), Gemini 3.8 Flash (Google DeepMind, 2026b) and Gemini 3.1 Pro preview (Google DeepMind, 2026a) through OpenRouter², whose weights cannot be hashed. Four input variants isolate the evidence: full, no graph, no retrieval, and no evidence. Each choice is scored by paired replay on the held-out grid: its run P95 in

each test block and its change relative to the control in the same block, with a Student- t interval over ten blocks.

Citations. For every valid full-context answer we record the kind of each cited item, whether a cited run belongs to the chosen or to a competing configuration, and its position in the evidence list. Since the schema admits only valid IDs, these counts describe which evidence a model presents.

Claims. Following error annotation in data-to-text evaluation (Thomson and Reiter, 2020), every quantitative or comparative claim in each distinct explanation is checked against the prompt evidence and labeled *correct*, *overstated*, *wrong* (including partly wrong), *wrong unit*, *wrong attribution* (a value assigned to the wrong source) or, for no-evidence prompts, *unsupported prior*. A claim is *misstated* if it is wrong, misattributed, overstated or in the wrong unit. A list-valued statement counts as one claim and a comparison as another; “consistently” or “always” lowest is overstated, since the chosen configuration is lowest in 8 of 10 retrieval blocks. Local models returned byte-identical text across seeds, so each distinct explanation is audited once per variant; hosted models varied across seeds and providers, so all distinct full-context explanations are audited. DeepSeek-V3.2’s providers ignored the 400-character limit in 14 of 20 answers; these are audited in full. An AI assistant drafted the labels under these rules; a script re-checked them where possible and one author verified all, changing none (Appendix C).

Reliance. Following the comprehensiveness test of ERASER (DeYoung et al., 2020), for each seed we delete the items cited in the full-context answer and, separately, the same number of uncited items of the same kinds. A decision that changes only under cited deletion indicates citation-specific reliance. A seed is non-estimable when no same-kind control exists, as for the training table, the only item of its kind.

5 Results

We report the audit at four levels, from the decision to the text that justifies it: what models choose (§5.1), what they cite (§5.2), what their explanations claim (§5.3) and whether the cited evidence matters (§5.4). A final subsection evaluates the testbed’s risk gate (§5.5).

¹<https://github.com/ggml-org/llama.cpp>, via llama-cpp-python 0.3.35.

²<https://openrouter.ai>

Table 2: LLM decisions with evidence (full, no graph, no retrieval), without it, and after deleting cited or matched uncited evidence; superscripts count seeds; n/e: non-estimable; [†]hosted. Paired test P95 change versus control at 5 and 8 ms (% , 95% half-width): ctrl 0; 50/1 -145 ± 125 and -180 ± 156 ; 1m/all -599 ± 188 and -631 ± 179 ; 1m/1 -631 ± 183 and -699 ± 206 .

Model	Evidence	None	Del. cited	Del. unc.
Qwen-7B	ctrl	50/1	n/e	n/e
Llama-8B	50/1	50/1	1m/all	ctrl
Qwen-14B	ctrl	50/1	n/e	n/e
Phi-4	ctrl	1m/1	50/1	ctrl
Llama-70B [†]	ctrl	1m/1	ctrl	ctrl
Qwen3-235B [†]	ctrl	1m/1 ³ ctrl	50/1	ctrl
DeepSeek-V3.2 [†]	ctrl	50/1	n/e	n/e
Nemotron-550B [†]	ctrl	ctrl	ctrl	ctrl
Gemini-Flash [†]	ctrl	50/1	n/e	n/e
Gemini-Pro [†]	ctrl	50/1	n/e	n/e

5.1 Decisions: Evidence Is Necessary and Sufficient

The table selector and its saturated-model equivalent choose the control, which is also best on the test grid, so any other choice is a loss; control test P95 is 179.7 ± 102.0 and $155.7 \pm 64.4 \mu\text{s}$ at 5 and 8 ms (95% half-widths across ten blocks). Matched-budget random search, Q-learning, UCB1 and BO converge on it as well, choosing it in 98–99% of 1,000 seeds at $B = 40$ (Appendix B).

Table 2 gives each model’s choice under every input condition. With evidence, nine of ten models choose the control in all three evidence variants, exactly as table lookup does; Llama-3.1-8B instead picks s50000_aone, which is worse on replay ($-145 \pm 125\%$ and $-180 \pm 156\%$). Removing the graph item changes no choice, and removing the retrieval runs changes none either: the training table alone suffices. Without evidence, every model except Nemotron, which declares the default, chooses a worse setting in at least four of five seeds, usually single-CPU affinity with $50 \mu\text{s}$ slack (six models) or with 1 ms slack (three). The evidence is therefore necessary; in this task it is also sufficient, since no model improves on the selector that simply reads it. 249 of 250 calls returned schema-valid output.

5.2 Citations: Valid, One-Sided and Front-Loaded

Every citation refers to a supplied item, since the schema permits nothing else, but what models choose to cite is far from a balanced account of the evidence. Across the 50 full-context answers, 285

Table 3: Verification of quantitative and comparative claims in explanations against the supplied evidence (AI-drafted labels, verified by an author). Corr.: correct; Over.: overstated; Wrong: wrong or partly wrong value; Unit/Src: wrong unit or evidence source; Prior: unsupported prior without evidence; Misst.: wrong, misattributed, overstated or wrong unit. Generated by claim_taxonomy.py.

Model	<i>n</i>	Corr.	Over.	Wrong	Unit/Src	Prior	Misst.
Qwen-7B	5	1	1	1	1	1	3
Llama-8B	7	–	–	6	–	1	5
Qwen-14B	6	1	1	–	3	1	4
Phi-4	10	5	3	1	–	1	4
<i>Local, all variants</i>	28	7	5	8	4	4	16
Llama-70B	7	5	–	2	–	–	2
Qwen3-235B	6	3	3	–	–	–	3
DeepSeek-V3.2	16	10	2	4	–	–	6
Nemotron-550B	23	13	5	5	–	–	10
Gemini-Flash	12	10	2	–	–	–	2
Gemini-Pro	5	3	2	–	–	–	2
<i>Hosted, full context</i>	69	44	14	11	–	–	25

citations comprise 25 of the training table, none of the interaction-graph item, and 260 of individual retrieval runs. Of these 260, 252 are runs of the configuration the model chose; only Gemini 3.1 Pro ever cites a run of a competing configuration. The cited evidence is thus confirmatory: no answer from the other nine models cites a run of a competing configuration, although s50000_aone beats the control in two of ten retrieval blocks.

Citations also favor the start of the evidence list. Retrieval runs appear in block order, and 209 of the 260 retrieval citations fall in the first five of ten blocks. Seven models cite exactly a prefix of the blocks, typically blocks 0–4 or 0–5 of the chosen configuration, and so never cite that configuration’s lowest run (block 7 for the control, block 8 for s50000_aone). The six-citation cap in the schema makes a prefix convenient, but it does not explain why the prefix is chosen over the most informative runs; the pattern is consistent with the position effects reported for long contexts (Liu et al., 2024).

5.3 Explanations: Cited but Not Correct

The explanations are often wrong about what the cited items say (Table 3): 16 of 28 audited claims from local models and 25 of 69 from hosted models are misstated. Errors fall into four recurring patterns (Table 4).

Overstated regularities. Thirteen of the 19 overstatements assert that the chosen configuration is lowest “consistently” or in “all” retrieval runs, whereas it is lowest in 8 of 10 retrieval blocks; the other six claim “significant” differences that

Table 4: Examples of misstated claims (paraphrased as audited) and the check against the supplied evidence; all latencies in μs .

Model, label	Claim	Evidence
Gemini-Pro [†] , overstated	control “consistently the lowest in retrieval runs”	lowest in 8 of 10 blocks
Nemotron [†] , wrong	s50000_aone ranges from 85.0 to 430.7	maximum is 571.84
DeepSeek [†] , wrong	control “consistently lower than any” s50000_aone run	ranges overlap: 85.033 < 183.217
Qwen-7B, wrong source	training table gives control 126.36	126.36 is one retrieval run; training mean 132.94
Qwen-14B, wrong unit	training table gives control 132.9413 ms	evidence is in μs
Qwen-14B, prior	“50000” denotes a smaller dataset size	no evidence given; name is opaque

no test in the evidence supports. These claims are easier to make when, as above, the cited runs are all of one configuration.

Range and extremum errors. All 11 wrong values from hosted models misreport a range, an extremum or a cross-configuration ordering. Nemotron repeatedly gives the s50000_aone range as 85.0–430.7 μs (its maximum is 571.84), and DeepSeek-V3.2 states that the control is “consistently lower than any” s50000_aone run although the ranges overlap (Figure 1). Llama-3.1-8B reports 571.84 μs , the *maximum* of s50000_aone, as the lowest P95, consistent with its wrong choice.

Source and unit confusion. Qwen2.5-7B attributes the retrieval value 126.36 to the training table, whose mean is 132.94; Qwen2.5-14B reports the training mean 132.9413 in milliseconds in three explanations, although the evidence declares microseconds.

Invented semantics. Without evidence, all four local models justify their choice by assigning meanings to opaque identifiers, for instance reading “50000” as a smaller dataset size or “one” as a more aggressive approach. The configuration names carry no such information.

These errors are largely decoupled from the decision: 35 of the 41 misstated claims accompany a choice of the control. Nor does variation in wording track variation in decisions: hosted models produced two to five distinct full-context explanations each while choosing the control in every seed. Valid citations thus establish provenance, and a correct decision can coexist with an incorrect account of the evidence that supports it.

5.4 Reliance: Rarely Estimable, Rarely Specific

Five models cite the training table in their full-context answers, so matched deletion is non-estimable for them. Among estimable models, only Phi-4 shows citation-specific reliance: deleting its six cited retrieval runs changes its choice in all five seeds while matched uncited deletion does not, but the new choice (s50000_aone) contradicts the training table that remains in the prompt. Llama-3.1-8B changes its choice under both deletions, which indicates sensitivity to context rather than reliance on what it cites. Llama-3.3-70B, Qwen3 and Nemotron change under neither, which is consistent with the remaining evidence still favoring the control but uninformative about reliance.

5.5 Enforcement: A Calibrated Gate under Shift

A test run is *unsafe* when more than 10% of its 40 events exceed a predeclared 200 μs deadline, a label computed from timing alone. For the gate, the score is a configuration’s training deadline-miss fraction, fixed before calibration on 4 ms runs only; all three prespecified $\alpha \in \{.05, .10, .20\}$ yield $\tau = .9225$ because the unsafe-score order statistics are tied. Every candidate is executed in an isolated process, so accepted and vetoed candidates both have observed outcomes (Figure 3 in Appendix B).

Unsafe acceptance is $5/25 = 20.0\%$ and $4/24 = 16.7\%$ among unsafe runs, above the nominal 5% and 10% levels; unsafe fractions among accepted runs are $5/20$ and $4/20$, and veto precision is $20/20$ in each period. As with LLM explanations, a nominally calibrated signal is only as good as the exchangeability it assumes, and a period shift is enough to break it. The bounded actuator restored both original settings in all 80 staged executions (Appendix B).

6 Discussion

Decisions and rationales come apart. Most mis-statements accompany a correct choice. This pattern is consistent with a decision driven by a simple reading of the evidence, here the minimum of the training table, and a rationale that is generated as a separate summary of the retrieval runs. Choosing the smallest of four means is easy; summarizing forty noisy measurements with ranges, extremes and quantifiers is where the errors concentrate. Because the decision is right, an end-to-end evaluation

would never surface these errors, and a reader who trusts the rationale would come away with a false picture of the evidence, for instance that the chosen configuration is never beaten. This mirrors findings that explanations can be plausible yet unfaithful to what drives a prediction (Jacovi and Goldberg, 2020; Turpin et al., 2023), with the difference that here the error is checkable against the prompt itself.

Citations as selective presentation. A citation list is also an editorial choice. Nine of ten models cited only runs of the configuration they chose, and most cited the first runs in the list, so a reader who follows the citations sees supporting evidence only and never the blocks in which a competitor wins. It also licenses overstatement (§5.3): “consistently the lowest” cannot be contradicted by the cited runs, only by those left uncited. The six-citation cap encourages selection, but not the systematic preference for confirming and early items.

Scale does not remove the error types. Every model we audited, including hosted frontier models, produced at least two misstated claims, and the same overstatement (“consistently the lowest”) recurs across model families. The local and hosted rates (16/28 and 25/69) are not directly comparable, since local explanations were audited across all input variants and hosted ones only with full context, but the error types are shared. The superlative, range and quantifier errors long documented in data-to-text generation (Thomson and Reiter, 2020; Chen et al., 2020a) are therefore not solved by instruction tuning or scale.

Implications for attribution metrics. Attribution benchmarks ask whether each cited passage supports its statement, often with an entailment judge (Rashkin et al., 2023; Gao et al., 2023b). One-sided citation lists are invisible to such per-citation scores, and many misstated claims quote correct numbers and err only in a comparison or quantifier, a near-miss that a coarse support judgment may accept; we did not run such judges. Evaluations of grounded rationales should therefore also measure what is left out, such as the share of contrary evidence cited, and check claims exactly where the evidence allows.

What deletion can and cannot show. Matched deletion measures whether a decision is sensitive to what is cited, which is a form of self-consistency rather than faithfulness to internal computation (Parcalabescu and Frank, 2024). Even that

is estimable only when the evidence contains redundant items of the same kind: half of our models cited the single training table, for which no control exists. Evidence sets for such audits should be designed with matched alternatives for every item a model might cite.

When could an LLM add value? Our evidence determines the best choice, so a selector that reads it is hard to beat. A language model could add value when the evidence is incomplete and prior knowledge, for example from documentation, carries information the measurements lack. Testing this requires tasks in which the default is not optimal and in which the baseline receives the same documents; the invented label meanings in our no-evidence runs show why such priors must be checked, not assumed.

7 Conclusion

A rationale that cites evidence promises a good decision, true claims about the evidence, and reliance on what it cites. We audited them separately for LLM configuration advice, with an equally informed baseline, evidence ablation and deletion, claim verification and independently labeled held-out outcomes. The promises came apart. Ten LLMs needed the evidence to choose well but added nothing beyond a selector that reads it; their citations were valid but one-sided (252 of 260 cited runs back the chosen configuration); 41 of 97 audited claims misstated validly cited evidence, mostly alongside a correct decision; and deletion showed citation-specific reliance for only one model.

For NLP evaluation, these results suggest three practices. First, citation validity and passage-level support are not enough: an explanation can cite valid items, omit the contrary ones and still misreport what it cites, so attribution should be checked claim by claim and for what is left out. Second, the errors that data-to-text research documented for tables, namely wrong numbers, superlatives, ranges and overstated regularities, persist in the free-text rationales of instruction-tuned models of every size we tested, and deserve targeted checks. Third, LLM decision support should be compared with a baseline given the same evidence, reporting decision quality and rationale correctness separately. Next steps are tasks where documentation priors could help, and multi-annotator or automatic claim checks.

Limitations

Task and generality. The audit covers one grounded-advice task with four candidate labels, English prompts and a small numeric evidence set that determines the best choice, and the default is best on the test grid. It can therefore expose missing added value and misreported evidence, but it cannot show that an LLM adds value on harder tasks, and its error rates need not transfer to domains with textual evidence (manuals, logs, documentation), longer contexts or other languages. The configuration labels are opaque by design; with descriptive labels, models’ priors could help or hurt differently.

Annotation. Claim labels were drafted with an AI assistant and verified by a single author, who changed none of them; this confirms the draft but is not an independent double annotation, so no inter-annotator agreement is reported, and the sample is small (97 claims). Segmenting explanations into claims and deciding when a claim is “overstated” involve judgment; we fixed both rules in advance (a list counts as one claim, a comparison as another; “consistently” is overstated when a configuration is lowest in 8 of 10 blocks), and supply every label with its check so that the annotation can be repeated or automated.

Models and prompting. All models share one system prompt and output schema at temperature 0; we did not vary prompt wording, evidence formatting, decoding or reasoning modes, each of which could change both decisions and explanations. Local explanations were byte-identical across seeds, and five greedy seeds are not independent samples. Local models were 4-bit quantized and may differ from full-precision weights. Hosted models were served by several providers, could not be hashed, varied across seeds and are not reproducible bit for bit; DeepSeek-V3.2’s providers ignored the explanation length limit. The six-citation cap and the 400-character limit shape what models cite and write, so the citation position effect in particular is confounded with the cap.

Reliance. Matched deletion measures the sensitivity of a decision to cited evidence, a form of self-consistency, not faithfulness to internal computation. It was non-estimable for half of the models, because they cited the only item of its kind.

Measurement scope. All measurements use one periodic workload family on one shared host, with

data roles separated by period rather than by workload family or hardware; run-level intervals do not account for uncontrolled interference or correlated system state and are unadjusted for multiple comparisons. Measured graph induction emits no edge, so the graph item is uninformative in this task. Selector, BO and RL comparisons replay recorded runs rather than executing closed loops, and because control is best they can reveal only losses. The telemetry prototype’s scheduler and block-I/O proxies are not end-to-end request latencies, and the local benchmarks in the appendix do not substitute for application workloads such as web serving, databases or GPU inference.

System and safety scope. The REST runtime is a dry run that never writes kernel settings, and the separate actuator changes only an owned child thread’s timer slack and affinity; neither establishes traffic isolation, a deployed multi-knob actuator or recovery from irreversible effects. The reasoner awaits model samples synchronously, and its logs are append-only files without tamper resistance. Prompting a model to respect graph edges does not enforce them; only the runtime’s explicit constraints are enforced. The rank gate’s assumptions fail under drift, selective labeling or a score updated with calibration data, as the period-shift result shows, and the gate replay neither verifies deployment assumptions nor establishes recovery. We therefore make no LLM-superiority, deployment-efficacy, no-worse-than-default or production-safety claim; Section 4 and Appendix A describe the comparisons and engineering such claims would require.

Ethics Statement

Intended use and risks of the findings. The paper studies how far users can trust LLM rationales that cite supplied evidence. Its main risk is misreading in either direction: treating valid citations or a correct final choice as evidence that a rationale is accurate, which our results argue against, or reading the audit as evidence that LLM configuration advisors are safe or effective, which it does not show (see Limitations). We recommend that LLM advice affecting real systems be paired with claim-level checks of its rationale and with enforcement that does not rely on the model.

Data, annotation and models. The data are timing measurements of local software workloads, syn-

thetic configuration labels, and model responses to one fixed technical prompt; they contain no personal data and involve no human participants. No annotators were recruited or paid; the explanation-claim labels were drafted with an AI assistant and verified by an author (see Use of AI assistants). All models were used for research inference under their licenses: Apache 2.0 (Qwen), the Llama 3.1 and 3.3 Community Licenses, MIT (Phi-4, DeepSeek-V3.2), OpenMDW-1.1 (Nemotron) and the Gemini API terms. Code, measurements, prompts, responses and annotations are provided as supplementary material and will be released publicly under Apache 2.0; the stored responses remain subject to each model’s terms.

System safety. The REST runtime performs no kernel writes; the experimental adapter changes only its owned child thread’s timer slack and affinity, restores both settings, and performs no host-wide writes. Broader interventions require an isolated testbed, documented rollback limits and appropriate authorization.

Compute. The study used about 3.2 CPU-hours of local inference without a GPU, 150 hosted API calls costing about 0.5 USD, and seconds of kernel measurement (Appendix C).

Use of AI assistants. The research idea, problem formulation and study design are the authors’. AI assistants were used for writing and debugging code, for drafting and restructuring prose, for locating references and for drafting the explanation-claim labels, which an author then verified; every reference was checked against its publisher, ACL Anthology, Crossref or arXiv record, and every claim label is released with its check so that it can be verified. The authors take full responsibility for the content. Every reported number is produced by the supplied scripts from raw logs whose SHA-256 hashes are recorded, so that it can be recomputed independently; no generated value substitutes for an unavailable observation.

References

- Marah Abdin, Jyoti Aneja, Harkirat Behl, et al. 2024. Phi-4 technical report. arXiv:2412.08905. <https://arxiv.org/abs/2412.08905>.
- Anastasios N. Angelopoulos, Stephen Bates, Adam Fisch, Lihua Lei, and Tal Schuster. 2024. Conformal

risk control. In *International Conference on Learning Representations*. <https://arxiv.org/abs/2208.02814>.

Pepa Atanasova, Oana-Maria Camburu, Christina Lioma, Thomas Lukasiewicz, Jakob Grue Simonsen, and Isabelle Augenstein. 2023. Faithfulness tests for natural language explanations. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, pages 283–294. <https://aclanthology.org/2023.acl-short.25/>.

Peter Auer, Nicolò Cesa-Bianchi, and Paul Fischer. 2002. Finite-time analysis of the multiarmed bandit problem. *Machine Learning*, 47(2–3):235–256. <https://doi.org/10.1023/A:1013689704352>.

Rina Foygel Barber, Emmanuel J. Candès, Aaditya Ramdas, and Ryan J. Tibshirani. 2023. Conformal prediction beyond exchangeability. *The Annals of Statistics*, 51(2):816–845. <https://doi.org/10.1214/23-AOS2276>.

Albert Bifet and Ricard Gavaldà. 2007. Learning from time-changing data with adaptive windowing. In *Proceedings of the 2007 SIAM International Conference on Data Mining*, pages 443–448. <https://doi.org/10.1137/1.9781611972771.42>.

Bernd Bohnet, Vinh Q. Tran, Pat Verga, Roei Aharoni, Daniel Andor, Livio Baldini Soares, Massimiliano Ciaramita, Jacob Eisenstein, et al. 2022. Attributed question answering: Evaluation and modeling for attributed large language models. arXiv:2212.08037. <https://arxiv.org/abs/2212.08037>.

Huilai Chen, Yuanbo Wen, Limin Cheng, Shouxu Kuang, et al. 2024. AutoOS: Make your OS more powerful by exploiting large language models. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *PMLR*, pages 7511–7525. <https://proceedings.mlr.press/v235/chen24at.html>.

Wenhu Chen, Jianshu Chen, Yu Su, Zhiyu Chen, and William Yang Wang. 2020a. Logical natural language generation from open-domain tables. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 7929–7942. <https://aclanthology.org/2020.acl-main.708/>.

Wenhu Chen, Hongmin Wang, Jianshu Chen, Yunkai Zhang, Hong Wang, Shiyang Li, Xiyu Zhou, and William Yang Wang. 2020b. TabFact: A large-scale dataset for table-based fact verification. In *International Conference on Learning Representations*. <https://arxiv.org/abs/1909.02164>.

Benjamin Cohen-Wang, Harshay Shah, Kristian Georgiev, and Aleksander Mądry. 2024. ContextCite: Attributing model generation to context. In *Advances in Neural Information Processing Systems*, volume 37. <https://arxiv.org/abs/2409.00729>.

880	DeepSeek-AI. 2025. DeepSeek-V3.2: Pushing the frontier of open large language models. arXiv:2512.02556. https://arxiv.org/abs/2512.02556 .	936
881		937
882		938
883		939
884	Jay DeYoung, Sarthak Jain, Nazneen Fatema Rajani, Eric Lehman, Caiming Xiong, Richard Socher, and Byron C. Wallace. 2020. ERASER: A benchmark to evaluate rationalized NLP models. In <i>Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics</i> , pages 4443–4458. https://aclanthology.org/2020.acl-main.408/ .	940
885		941
886		942
887		
888		943
889		944
890		945
891	Jesse Dodge, Suchin Gururangan, Dallas Card, Roy Schwartz, and Noah A. Smith. 2019. Show your work: Improved reporting of experimental results. In <i>Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)</i> , pages 2185–2194. https://aclanthology.org/D19-1224/ .	946
892		947
893		948
894		
895		949
896		950
897		951
898		952
		953
899	António Farinhas, Chrysoula Zerva, Dennis Ulmer, and André F. T. Martins. 2024. Non-exchangeable conformal risk control. In <i>International Conference on Learning Representations</i> . https://arxiv.org/abs/2310.01262 .	954
900		955
901		956
902		957
903		
904	Luyu Gao, Zhu Yun Dai, Panupong Pasupat, Anthony Chen, Arun Tejasvi Chaganty, Yicheng Fan, Vincent Y. Zhao, Ni Lao, Hongrae Lee, Da-Cheng Juan, and Kelvin Guu. 2023a. RARR: Researching and revising what language models say, using language models. In <i>Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)</i> , pages 16477–16508. https://aclanthology.org/2023.acl-long.910/ .	958
905		959
906		960
907		961
908		962
909		
910		963
911		964
912		965
913		966
914		967
915		968
916		
917		969
918		970
919		971
920		972
921		973
922		974
923		
924		975
925		976
926		977
927		978
928		979
929		980
930	Yonatan Geifman and Ran El-Yaniv. 2017. Selective classification for deep neural networks. In <i>Advances in Neural Information Processing Systems</i> , volume 30, pages 4878–4887. https://papers.nips.cc/paper/7073-selective-classification-for-deep-neural-networks .	981
931		982
932		983
933		984
934		985
935		
		986
		987
		988
		989
		990
		991
		992
		993
		994
		995
		996
		997
		998
		999
		1000

- Jiale Lao, Yibo Wang, Yufei Li, Jianping Wang, Yunjia Zhang, Zhiyuan Cheng, Wanghu Chen, Mingjie Tang, and Jianguo Wang. 2024. GPTuner: A manual-reading database tuning system via GPT-guided Bayesian optimization. *Proceedings of the VLDB Endowment*, 17(8):1939–1952. <https://doi.org/10.14778/3659437.3659449>.
- Hyeonmyeong Lee, Sungmin Jung, and Heeseung Jo. 2022. STUN: Reinforcement-learning-based optimization of kernel scheduler parameters for static workload performance. *Applied Sciences*, 12(14):7072. <https://doi.org/10.3390/app12147072>.
- Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, et al. 2020. Retrieval-augmented generation for knowledge-intensive NLP tasks. In *Advances in Neural Information Processing Systems*, volume 33, pages 9459–9474. <https://proceedings.neurips.cc/paper/2020/hash/6b493230205f780e1bc26945df7481e5-Abstract.html>.
- Hongyu Lin, Yuchen Li, Haoran Luo, et al. 2025. TuneAgent (first released as OS-R1): Agentic operating system kernel tuning with reinforcement learning. arXiv:2508.12551. <https://arxiv.org/abs/2508.12551>.
- Linux man-pages project. 2026a. PR_SET_TIMERSLACK(2const). https://man7.org/linux/man-pages/man2/PR_SET_TIMERSLACK.2const.html.
- Linux man-pages project. 2026b. sched_setaffinity(2). https://man7.org/linux/man-pages/man2/sched_setaffinity.2.html.
- Nelson F. Liu, Kevin Lin, John Hewitt, Ashwin Paranjape, Michele Bevilacqua, Fabio Petroni, and Percy Liang. 2024. Lost in the middle: How language models use long contexts. *Transactions of the Association for Computational Linguistics*, 12:157–173. <https://aclanthology.org/2024.tacl-1.9/>.
- Nelson F. Liu, Tianyi Zhang, and Percy Liang. 2023. Evaluating verifiability in generative search engines. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 7001–7025. <https://aclanthology.org/2023.findings-emnlp.467/>.
- Shayne Longpre, Kartik Perisetla, Anthony Chen, Nikhil Ramesh, Chris DuBois, and Sameer Singh. 2021. Entity-based knowledge conflicts in question answering. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 7052–7063. <https://aclanthology.org/2021.emnlp-main.565/>.
- Andreas Madsen, Sarath Chandar, and Siva Reddy. 2024. Are self-explanations from large language models faithful? In *Findings of the Association for Computational Linguistics: ACL 2024*, pages 295–337. <https://aclanthology.org/2024.findings-acl.19/>.
- Joshua Maynez, Shashi Narayan, Bernd Bohnet, and Ryan McDonald. 2020. On faithfulness and factuality in abstractive summarization. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 1906–1919. <https://aclanthology.org/2020.acl-main.173/>.
- Sewon Min, Kalpesh Krishna, Xinxu Lyu, Mike Lewis, Wen-tau Yih, Pang Wei Koh, Mohit Iyyer, Luke Zettlemoyer, and Hannaneh Hajishirzi. 2023. FActScore: Fine-grained atomic evaluation of factual precision in long form text generation. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*. <https://arxiv.org/abs/2305.14251>.
- Christopher Mohri and Tatsunori Hashimoto. 2024. Language models with conformal factuality guarantees. In *Proceedings of the 41st International Conference on Machine Learning*. <https://arxiv.org/abs/2402.10978>.
- NVIDIA. 2026. Nemotron 3 Ultra: Open, efficient mixture-of-experts hybrid Mamba-transformer model for agentic reasoning. Technical report. <https://research.nvidia.com/labs/nemotron/files/NVIDIA-Nemotron-3-Ultra-Technical-Report.pdf>.
- Letitia Parcalabescu and Anette Frank. 2024. On measuring faithfulness or self-consistency of natural language explanations. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 6048–6089. <https://aclanthology.org/2024.acl-1.329/>.
- Victor Quach, Adam Fisch, Tal Schuster, Adam Yala, Jae Ho Sohn, Tommi S. Jaakkola, and Regina Barzilay. 2024. Conformal language modeling. In *International Conference on Learning Representations*. <https://arxiv.org/abs/2306.10193>.
- Qwen Team. 2024. Qwen2.5 technical report. arXiv:2412.15115. <https://arxiv.org/abs/2412.15115>.
- Hannah Rashkin, Vitaly Nikolaev, Matthew Lamm, Lora Aroyo, Michael Collins, Dipanjan Das, Slav Petrov, Gaurav Singh Tomar, Iulia Turc, and David Reitter. 2023. Measuring attribution in natural language generation models. *Computational Linguistics*, 49(4). <https://arxiv.org/abs/2112.12870>.
- Freda Shi, Xinyun Chen, Kanishka Misra, Nathan Scales, David Dohan, Ed Chi, Nathanael Schärli, and Denny Zhou. 2023. Large language models can be easily distracted by irrelevant context. In *Proceedings of the 40th International Conference on Machine Learning*. <https://arxiv.org/abs/2302.00093>.
- Jasper Snoek, Hugo Larochelle, and Ryan P. Adams. 2012. Practical Bayesian optimization of machine learning algorithms. In *Advances in Neural Information Processing Systems*, volume 25, pages 2951–2959. <https://proceedings.neurips.cc/paper>

1106	/2012/hash/05311655a15b75fab86956663e1819cd-Abstract.html .	Revealing the behavior of large language models in knowledge conflicts. In <i>International Conference on Learning Representations</i> . https://arxiv.org/abs/2305.13300 .	1163
1107			1164
1108	Anders Søgaard, Sebastian Ebert, Jasmijn Bastings, and Katja Filippova. 2021. We need to talk about random splits. In <i>Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics: Main Volume</i> , pages 1823–1832. https://aclanthology.org/2021.eacl-main.156/ .		1165
1109			1166
1110		Miao Xiong, Zhiyuan Hu, Xinyang Lu, Yifei Li, Jie Fu, Junxian He, and Bryan Hooi. 2024. Can LLMs express their uncertainty? an empirical evaluation of confidence elicitation in LLMs. In <i>International Conference on Learning Representations</i> . https://arxiv.org/abs/2306.13063 .	1167
1111			1168
1112			1169
1113			1170
1114			1171
1115	Craig Thomson and Ehud Reiter. 2020. A gold standard methodology for evaluating accuracy in data-to-text systems. In <i>Proceedings of the 13th International Conference on Natural Language Generation</i> , pages 158–168. https://aclanthology.org/2020.inlg-1.22/ .		1172
1116		An Yang et al. 2025. Qwen3 technical report. arXiv:2505.09388. https://arxiv.org/abs/2505.09388 .	1173
1117			1174
1118			1175
1119		Chengrun Yang, Xuezhi Wang, Yifeng Lu, Hanxiao Liu, Quoc V. Le, Denny Zhou, and Xinyun Chen. 2024. Large language models as optimizers. In <i>International Conference on Learning Representations</i> . https://arxiv.org/abs/2309.03409 .	1176
1120			1177
1121	Katherine Tian, Eric Mitchell, Allan Zhou, Archit Sharma, Rafael Rafailov, Huaxiu Yao, Chelsea Finn, and Christopher D. Manning. 2023. Just ask for calibration: Strategies for eliciting calibrated confidence scores from language models fine-tuned with human feedback. In <i>Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing</i> , pages 5433–5442. https://aclanthology.org/2023.emnlp-main.330/ .		1178
1122			1179
1123			1180
1124		Ji Zhang, Yu Liu, Ke Zhou, Guoliang Li, et al. 2019. An end-to-end automatic cloud database tuning system using deep reinforcement learning. In <i>Proceedings of the 2019 International Conference on Management of Data (SIGMOD)</i> , pages 415–432. https://doi.org/10.1145/3299869.3300085 .	1181
1125			1182
1126			1183
1127			1184
1128			1185
1129			1186
1130	Immanuel Trummer. 2022. DB-BERT: A database tuning tool that “reads the manual”. In <i>Proceedings of the 2022 International Conference on Management of Data (SIGMOD)</i> . https://doi.org/10.1145/3514221.3517843 .	A GroundKern Testbed Details	1187
1131			
1132		Design hypotheses. GroundKern is motivated by three hypotheses. First, a configuration’s utility may depend on workload and hardware. Second, joint knob effects may be non-additive. Third, performance, failure probability and intervention cost may conflict. Each hypothesis requires measured interventions with explicit baselines.	1188
1133			1189
1134			1190
1135	Miles Turpin, Julian Michael, Ethan Perez, and Samuel R. Bowman. 2023. Language models don’t always say what they think: Unfaithful explanations in chain-of-thought prompting. In <i>Advances in Neural Information Processing Systems</i> . https://arxiv.org/abs/2305.04388 .		1191
1136			1192
1137			1193
1138			1194
1139		For a loss L , define an interaction contrast from four independently measured configurations:	1195
1140			1196
1141	Dana Van Aken, Andrew Pavlo, Geoffrey J. Gordon, and Bohan Zhang. 2017. Automatic database management system tuning through large-scale machine learning. In <i>Proceedings of the 2017 ACM International Conference on Management of Data (SIGMOD)</i> , pages 1009–1024. https://doi.org/10.1145/3035918.3064029 .	$I_{ab} = L(a, b) - L(a, b_0) - L(a_0, b) + L(a_0, b_0).$	1197
1142			
1143		For lower-is-better loss, negative I_{ab} indicates a beneficial non-additive interaction relative to the chosen baseline. A bootstrap interval must resample independent runs, not individual requests from a single run. The sign can change with workload, value range or kernel version, so an edge needs its scope, configuration values, measurement IDs and uncertainty interval. Such an edge describes an empirical interaction, not an unconditional causal coupling.	1198
1144			1199
1145			1200
1146			1201
1147			1202
1148	Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc V. Le, Ed H. Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. 2023. Self-consistency improves chain of thought reasoning in language models. In <i>International Conference on Learning Representations</i> . https://arxiv.org/abs/2203.11171 .		1203
1149			1204
1150			1205
1151			1206
1152			1207
1153		A DEPENDS-ON edge requires a validated admissibility rule or measured range conditioning; it cannot be inferred solely from the sign of I_{ab} .	1208
1154			1209
1155	Sam Wiseman, Stuart M. Shieber, and Alexander M. Rush. 2017. Challenges in data-to-document generation. In <i>Proceedings of the 2017 Conference on Empirical Methods in Natural Language Processing</i> , pages 2253–2263. https://aclanthology.org/D17-1239/ .		1210
1156		The seed graph is empty until edges pass this test.	1211
1157		The principal question is whether a reasoner can	1212
1158			
1159			
1160			
1161	Jian Xie, Kai Zhang, Jiangjie Chen, Renze Lou, and Yu Su. 2024. Adaptive chameleon or stubborn sloth:		
1162			

Algorithm 1 Implemented dry-run bundle path

Require: Validated bundle B , fixed score function, calibration threshold τ

- 1: Validate every member and joint constraints before changing state
- 2: **if** another bundle is active **then**
- 3: Return conflict
- 4: **end if**
- 5: **if** $\max_{a \in B} u(a) \geq \tau$ **then**
- 6: Record veto for all members; return
- 7: **end if**
- 8: Stage the complete bundle in memory at 5%
- 9: **for** requested stage in $\{25, 50, 100\}\%$ **do**
- 10: Advance the simulated bundle; record the transition
- 11: **end for**
- 12: On cancellation, clear the complete simulated bundle
- 13: Never infer an unsafe/safe training label from simulated completion

use independently acquired evidence more effectively than an optimizer with the same inputs, after charging the cost of acquiring that evidence.

Disagreement score. For a knob–value pair proposed in v of k samples, the reasoner’s score is $u = 0.7(1 - v/k) + 0.3(1 - w)$, clipped to $[0.02, 1]$, where w is the strongest retrieved synergy weight (zero when absent). It is neither predictive entropy nor a calibrated probability. Failed samples stay in the denominator, and repeated copies of a fallback are not counted as independent votes.

Bundle contract. A request contains a nonempty bundle identifier and a list of recommendations with id, knob, proposed, uncertainty, and optional rationale and provenance. Each request submits one bundle. Duplicate knobs, nonfinite scores, unsupported controls, and example-policy bound violations reject the entire request. Dirty-ratio changes must supply both values and satisfy the example policy that the background ratio is less than the foreground ratio. This is a declared policy, not proof that a configuration improves any workload. If any score satisfies $u \geq \tau$, all members are vetoed; otherwise all are staged together. A concurrent bundle receives a conflict response. Dry-run results return `applied=[]` and list simulated members separately.

Deployment requirements. A real actuator must snapshot all old values, validate the full target configuration and available host interfaces, serialize updates, verify readback, and restore every old value on failure. Multiple `sysctl` writes are not a kernel transaction: partial exposure must be measured and cannot be described as atomic hardware state. Canary shares require separate hosts or an appropriate isolation mechanism; changing a host-global `sysctl` cannot isolate a fraction of requests on that host. Missing, stale or malformed end-to-end telemetry must stop promotion. These capabilities are requirements, not features demonstrated by the dry-run state machine.

One illustrative decision trace. Consider supplied context with example values 5 and 15 for the two dirty ratios, and retrieved evidence IDs e_1, t_1 . Suppose all three model samples agree on this pair and $w = 0$, giving $u = 0.30$. If an independently obtained threshold is $\tau = 0.25$, the whole bundle is vetoed; if $\tau = 0.35$, it can enter the dry-run canary after joint validation. The audit record includes the values, scores, threshold and decision. These numbers illustrate the contract and are not a measured recommendation or calibrated operating point. The initial artifact threshold is zero until calibration data is explicitly supplied.

Actuator. A separate adapter snapshots a child thread’s timer slack and CPU affinity ([Linux man-pages project, 2026a,b](#)), validates the target, applies both settings, verifies readback and restores both originals in a finally path, including after partial application. Updates are sequential, not an atomic kernel transaction.

Expected cost of gated acceptance. Let $Y = 1$ denote an independently defined unsafe outcome and $A = 1$ acceptance. An outcome label is not defined by whether a rollback happened. For example, a benchmark can predeclare a request deadline, count observed failures or deadline exceedances, and define an action-level breach from these raw observations in a fixed window. Rejected actions have unknown counterfactual outcomes unless evaluated in an isolated experiment; they must not automatically be labeled safe or unsafe.

Theorem 1 (Cost identity). *For costs $c_{FN}, c_{FP}, \lambda \geq 0$, define $C = c_{FN}\mathbf{1}\{Y = 1, A = 1\} + c_{FP}\mathbf{1}\{Y = 0, A = 0\} + \lambda D$, where $D \geq 0$ is measured exposure or SLO debt. Write $\pi = \Pr(Y = 1)$, $q = \Pr(A = 1 \mid Y = 1)$ and*

$r = \Pr(A = 0 \mid Y = 0)$. Then

$$\mathbb{E}C = c_{\text{FN}}\pi q + c_{\text{FP}}(1 - \pi)r + \lambda \mathbb{E}D.$$

For a zero-probability class its associated product is zero.

Proof. Expand each indicator expectation as its joint probability and use the definition of conditional probability. Linearity requires no independence. $\square$

This is an accounting identity, not a claim of novelty or a safety bound. A bound $q \leq \alpha$ controls misses among unsafe actions. A different bound $\Pr(Y = 1 \mid A = 1) \leq \alpha$ instead gives $\Pr(Y = 1, A = 1) \leq \alpha \Pr(A = 1)$, not $\pi\alpha$. Likewise r is a veto rate among safe actions, not the proportion of vetoed actions that are safe. We report all denominators separately.

For bundle-level calibration, u denotes the maximum member score and Y the independently observed bundle outcome. Per-knob labels alone do not establish a bundle-level guarantee.

Selecting a smaller threshold by a cost rule cannot enlarge the acceptance event. Selecting a larger threshold can invalidate the bound; therefore the implementation uses the minimum of the rank threshold and optional cost-selected cap. No unsafe calibration examples implies veto-all. Calibration-window diagnostics are not test-set estimates.

Drift and selective observation. Sliding windows and a drift detector (Bifet and Gavaldà, 2007) do not guarantee exchangeability or a deterministic recovery time. Under an explicitly justified total-variation bound on the *joint calibration–test law*, the probability of the same event changes by at most that bound; a marginal score-distance statement alone does not establish it. We neither estimate that distance nor claim recovery. Online labels selected by acceptance introduce additional selection bias. The runtime holds its threshold at zero on detected change until an explicit refresh; refresh remains an exploratory operation, not proof of safety.

Protocol requirements beyond this study. Workload, hardware and session groups are fixed before graph induction and assigned to disjoint train, retrieval, calibration and test roles; hashes of graphs, sources, prompts and calibrators are frozen before test results are opened, and a validator rejects overlap or unknown lineage. For

each bundle we count proposed, rejected, vetoed, accepted, staged, completed and rolled-back outcomes, report gate TP/FP/FN/TN with both $\#$ unsafe-accepted/ $\#$ unsafe and $\#$ accepted, and leave unknown labels unknown; rollback precision needs independently labeled executed actions. A lower threshold accepts a subset of fixed candidates, but the unsafe fraction among survivors need not be monotone. Absolute per-workload metrics precede ratios against a named baseline; a Pareto claim needs several measured operating points per method. Drift studies must keep the threshold trajectory and held-out outcomes at each decision, and faithfulness is tested by deleting cited and matched uncited context under fixed seeds.

B Additional Measurements

B.1 Selector Replay on the Follow-up

All controlled runs used one Intel Core i5-3570 host (four logical CPUs, Linux 6.17.0-23-generic). Each fresh child set its timer slack with PR_SET_TIMERSLACK and its affinity with sched_setaffinity, checked both by read-back, and recorded 40 wake-ups against absolute monotonic-clock targets after five warm-ups.

The empirical selector minimizes mean training-run P95 over the four configurations. A saturated two-factor representation encodes the same four means as a baseline, two main effects and their interaction; minimizing it adds no information to table lookup. Both selectors have the same 40 training runs and zero online search trials. An unchanged-control selector and a seeded uniform-random selector provide references. Choices are frozen before calibration and test. We evaluate these fixed selectors by paired *offline replay* on the same measured test grid, not by separate closed-loop executions. A block is the uncertainty unit, and paired reductions are calculated within each period against its explicit control.

Both informed selectors chose 50 μ s slack and all CPUs, identical to control; their paired improvement is exactly zero in this replay. Control P95 is 179.7 ± 102.0 and $155.7 \pm 64.4 \mu$ s at 5 and 8 ms (95% Student- t half-width across ten blocks). The training interaction contrast is -18.6μ s; the pre-specified normalized-contrast bootstrap emitted no edge because its interval included zero. We do not turn a nonsignificant contrast into a dependency rule or a semantic-reasoning gain. Acquisition took 6.54 s for training, 8.29 s for retrieval and 10.15 s

Table 5: Follow-up selector replay: mean of ten test-run P95 values. Misses are observed lateness above 200 μ s. Shared rows across selectors are intentional: selectors choosing the same configuration use the same held-out observations.

Period	Selector	P95 (μ s)	Misses
5	control	179.7	13/400
5	empirical min	179.7	13/400
5	graph min	179.7	13/400
5	seeded random	811.6	248/400
8	control	155.7	13/400
8	empirical min	155.7	13/400
8	graph min	155.7	13/400
8	seeded random	642.2	225/400

Table 6: Matched-budget replay: % of 1,000 seeds choosing control/expected paired P95 change at 5 ms (%); last column: mean training evaluations of 1 ms settings at $B = 40$.

Method	$B = 4$	$B = 8$	$B = 16$	$B = 40$	1 ms evals
Random search	62/-55	74/-38	85/-22	98/-3	19.8
ϵ -greedy (RL)	65/-51	86/-20	95/-8	99/-1	3.8
UCB1 (RL)	65/-51	77/-33	91/-12	99/-1	12.0
GP-EI (BO)	65/-51	83/-25	95/-8	99/-1	2.5

for calibration; graph construction and selection took 0.10 s, included in calibration-phase wall time.

B.2 Matched-Budget Optimizers

Random search, ϵ -greedy Q-learning ($\epsilon = 0.1$), UCB1 and GP-EI BO each evaluate $B \in \{4, 8, 16, 40\}$ recorded training runs, starting with one per configuration, and recommend their lowest observed or posterior mean, scored by the same replay over 1,000 seeds (Table 6). Control is best at both test periods, so this comparison can reveal only losses. With small budgets, noisy single-CPU runs mislead the methods; BO spends fewest evaluations on the 1 ms settings, whose training runs miss over 90% of deadlines.

B.3 Gate Replay

B.4 Staged Actuation

Across 80 fresh executions, 66 stop on a stage breach and 14 complete all stages; all 80 restore and read back both original values. At 5 and 8 ms, observed deadline misses are respectively 186/816 and 190/880; summed exposure durations are 4.106 and 7.066 s. Early stops reduce the number of observed events, so these fractions cannot be compared directly with fixed-duration, no-stop runs to claim an anomaly reduction. The stop criterion itself is not independent ground truth




Figure 3: Frozen-gate replay on the 80 held-out test runs; each point is one run. Positive means unsafe (dashed line) and a gate-positive decision is veto. Counts at every tested α (all give $\tau = .9225$): 5 ms, 20 accepted / 20 vetoed, TP 20, FP 0, FN 5, TN 15; 8 ms, 20 / 20, TP 20, FP 0, FN 4, TN 16. Blue points above the line are the accepted unsafe runs. Generated from raw logs by `plot_gate_replay.py`.

Table 7: Measured local baseline. Mean of run medians, and mean of run P95 values with 95% Student- t CI half-width across 10 runs. Last column is the percentage of errors or operations exceeding a predeclared 50 ms deadline. No tuning gain is implied.

Workload	Median ms	P95 ms	Events %
file_fsync	1.903	2.035 ± 0.043	0.00
sqlite_commit	1.870	2.187 ± 0.273	0.00
zlib_compress	1.474	2.038 ± 0.356	0.00

for rollback precision.

B.5 Local Baseline Measurements

We executed three bounded microbenchmarks on the available host: an Intel Core i5-3570 with four logical CPUs, approximately 31.2 GiB RAM, and Linux 6.17.0-23-generic. These runs characterize the measurement pipeline: no kernel knob was modified and no LLM or controller was compared.

Each workload had 10 separate invocations of 100 measured operations, preceded by 10 warm-up operations. Workload-run order was shuffled with seed 4088. SQLite inserts a 4 KiB value and commits each transaction in WAL mode with synchronous=FULL. The file test overwrites a 64 KiB region and calls fsync. The compression test compresses a fixed seeded 64 KiB payload at zlib level 6 and verifies decompression. Latency is wall-clock operation duration from a monotonic nanosecond clock; the compression measurement includes verification. The raw log retains every operation, run ID, deadline and error flag, with a SHA-256 hash and software/environment manifest.

Table 8: Held-out kernel intervention measurements. P95 is the mean of ten run-level P95 wake-up latenesses. Misses count lateness exceeding the predeclared 200 μ s deadline. CPU affinity is all four CPUs or one CPU.

Period	Slack/CPU	P95 (μ s)	Misses
2	50us/all	89.5	2/800
2	50us/one	109.4	18/800
2	1ms/all	1016.8	800/800
2	1ms/one	968.9	771/800
5	50us/all	94.2	2/800
5	50us/one	133.8	22/800
5	1ms/all	1022.6	663/800
5	1ms/one	1073.7	680/800

The anomaly denominator is all measured operations, including failures; an operation counts once if either it fails or exceeds the deadline. This label is independent of veto and rollback. Zero events in these short runs is not evidence of zero population risk. We do not pool heterogeneous workload latencies or infer kernel tuning effects from these baseline timings.

B.6 Temporal-Holdout Kernel Interventions

We additionally executed real Linux timer-slack and CPU-affinity interventions in fresh owned child processes, without modifying host-wide sysctls. The two factors were timer slack (50 μ s or 1 ms, set by PR_SET_TIMERSLACK) and affinity (all four available CPUs or CPU 0). Every child checked kernel readback before recording events. This isolated measurement driver is separate from the dry-run GroundKern actuator.

For each periodic wake-up interval (2 or 5 ms), five training blocks and ten subsequent test blocks evaluated all four configurations. Within each block, the eight interval/configuration cases were randomly ordered using seed 4088. Each fresh process performed ten warm-up and 80 measured wake-ups against absolute monotonic-clock targets. The plan was written before execution; the configuration minimizing mean training-run P95 lateness was frozen before test execution. This is temporal holdout within one host and workload family, not the independent workload/hardware split required for claims of generalization. The complete 120-process experiment retained 9,600 event timestamps, readbacks, frozen selections and hashes. No external competing load was introduced.

The training selector chose the explicit control (50 μ s slack, all CPUs) at both intervals: there is no selected-policy improvement over that control.

Its test P95 was 89.5 ± 9.2 and $94.2 \pm 6.9 \mu$ s at 2 and 5 ms, respectively (95% Student-*t* half-width across ten runs). Raising slack to 1 ms with all CPUs increased P95 by 927.3 ± 16.4 and $928.4 \pm 10.0 \mu$ s in paired blocks. Affinity-only differences were 19.9 ± 36.7 and $39.6 \pm 62.6 \mu$ s; these intervals include zero. The slack effect is consistent with timer coalescing; no power measurement, graph discovery, LLM advantage or staged rollback is inferred. A miss is defined by observed wake-up time, independently of any safety gate.

C Prompt and Output Schema

Every model receives the following system prompt:

Select one allowed config to minimize periodic wake-up P95 latency. Return JSON with config (exact allowed name), cited_ids (evidence IDs), and explanation. Evidence contains measured training summaries and retrieval runs, not held-out test outcomes. Do not invent evidence.

The user message is a JSON object with the four allowed configuration names and the evidence list for the input variant. Items have the forms

```
{
  "id": "training-table", "kind": "table",
  "values": {
    "s50000_aall": 132.9413, ...
  }
}, {
  "id": "interaction", "kind": "graph",
  "contrast_us": -18.62, "edges": [],
  "scope": "same measured table; graph is a reparameterization, not extra evidence"
}, {
  "id": "retrieval-3-0-s50000_aall",
  "kind": "retrieval", "config": "s50000_aall",
  "period_ns": 3000000, "p95_us": 126.36
}
```

The output schema requires config from the allowed list, cited_ids with at most six entries restricted to the supplied IDs (none for the no-evidence variant), and an explanation of at most 400 characters. Local models enforce it as a llama.cpp JSON grammar with greedy decoding; hosted models receive it as a strict JSON schema through OpenRouter, restricted to providers that honor it. Every answer is re-validated, and an invented or deleted citation makes it invalid; the explanation length is not re-checked, and 14 of 20 DeepSeek-V3.2 answers exceed 400 characters. Claim and citation counts in §5 are recomputed from these records by claim_taxonomy.py and citation_analysis.py. A single script, reproduce_paper.py, reruns every analysis on a copy of the raw logs and checks that each table and quoted number in the paper is reproduced exactly; the supplementary ARTIFACTS.md maps each result

to its data and script and includes a data statement. The supplementary records contain each prompt, its SHA-256 hash, the raw response and, for hosted models, the routed provider.

Claim-label checks and intervals. The 97 claim labels were drafted with an AI assistant. A rule-based check (`claim_label_autocheck.py`) re-derives the expected label of 56 claims directly from the evidence (values attributed to the training table, lists of runs, superlatives, ranges, thresholds, units and no-evidence claims) and agrees with every drafted label; flipping each label between *correct* and *wrong* is detected in 76 of 97 cases. The remaining 41 claims require judgment. One author then reviewed every label against the evidence, reading these 41 claims in full, and changed none (agreement with the draft 100%); because the reviewer checked an existing draft rather than labeling independently, this is a confirmation, not an inter-annotator agreement. With 95% Wilson intervals, the misstated proportions are 57% (39–73%) for local models, 36% (26–48%) for hosted models and 42% (33–52%) overall; 97% (94–98%) of cited retrieval runs belong to the chosen configuration and 80% (75–85%) come from the first half of the evidence list. These intervals treat claims and citations as independent, which overstates their precision because hosted explanations repeat claims across seeds.

Compute budget. Local models ran as 4-bit GGUF files (Q4_K_M; Phi-4 Q4_K) with llama-cpp-python 0.3.35 (<https://github.com/abetlen/llama-cpp-python>) on a four-core CPU container without a GPU, with an 8,192-token context and at most 384 output tokens. Their 100 calls took 11,580 s of wall time (about 3.2 CPU-hours): Qwen2.5-7B 1,603 s, Llama-3.1-8B 2,325 s, Qwen2.5-14B 3,136 s and Phi-4 4,516 s. The 150 hosted calls through OpenRouter, with at most 4,096 output tokens, took 1,224 s in total and cost about 0.5 USD. Kernel measurements ran on the Intel Core i5-3570 host described in Appendix B; acquiring the training, retrieval and calibration roles took 6.54, 8.29 and 10.15 s. Optimizer replays over 1,000 seeds take seconds on one CPU.